Experiments continued

What matters for compatibility recovery?

Stephanus Park, 15.05.26 — Pitch

Stephanus Park, 15.05.26 — Pitch

Recap

Where we stopped last time

Key points

- Metadata routing changes expert selection but does not reliably improve utility
- Learned proxy models and metadata-conditioned variants did not recover utility-aligned rankings
- The bottleneck was localized to compatibility estimation, not the routing mechanism itself

The new question
How can we estimate compatibility without leakage?

I tested three directions:

- Learned utility predictors
- Risk-constrained learned response routing
- Target-local support-set NELBO

Results

Stephanus Park, 15.05.26 — Pitch

Result A

Learned utility is not yet adoption-ready

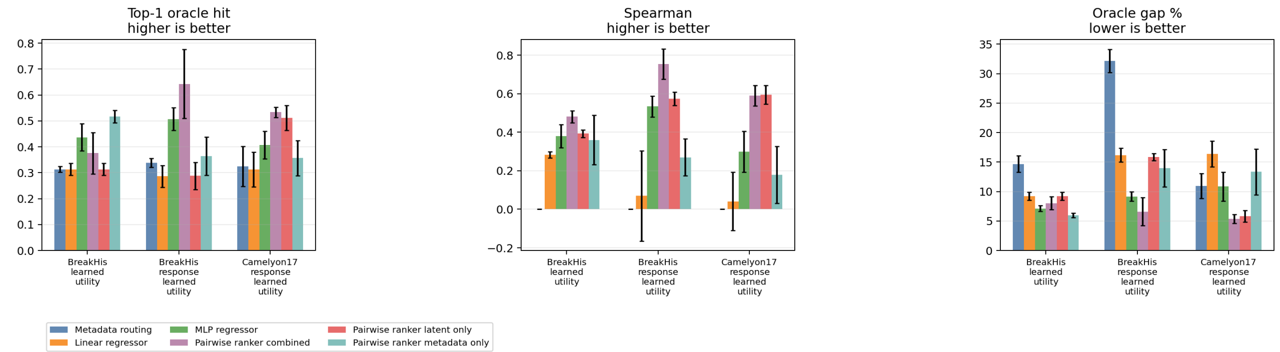

Stephanus Park, 15.05.26 — Pitch

Result B

Risk-constrained response routing is a weak pass

Result C

Support-Set NELBO is a strong positive result